

Sakshi Shukla

Data Engineer

+91-9696404772 | sakshishukla.tech@gmail.com | [LinkedIn](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Data Engineer with **2+ years** of experience building data pipelines, ETL/ELT workflows, and cloud-based data solutions. Proficient in **Python, SQL, PySpark, Databricks, Azure, AWS, Apache Airflow, and Apache Kafka**. Experienced in designing scalable data platforms, implementing data quality frameworks, Medallion Lakehouse Architecture, dimensional data modeling, and reliable analytics solutions.

SKILLS

Languages: Python, SQL

Big Data & Processing: Apache Spark (PySpark), Databricks, Delta Lake, Hive, Apache Kafka (CDC), Apache Airflow

Cloud Platforms: AWS (S3, Glue, Lambda, Redshift, EventBridge, SQS, SNS), Azure (Azure Databricks, ADLS Gen2)

Data Modeling & DWH: Dimensional Modeling, Star Schema, SCD Type 2, Medallion Lakehouse (Bronze / Silver / Gold), ETL / ELT, Data Quality Frameworks

Databases & Tools: PostgreSQL, Redshift, Hive Metastore, Git, Jira, REST APIs

EXPERIENCE

Data Engineer

June 2024 – Present

Tata Consultancy Services (TCS)

India

- Designed and developed ingestion pipelines integrating data from 5+ enterprise source systems, including POS, CRM, ERP, and e-commerce platforms using batch processing, CDC (Kafka), and REST APIs, reducing reporting time from 3-4 days to near real-time.
- Built scalable ETL pipelines in PySpark to cleanse, standardize, and transform data across Bronze, Silver, and Gold layers following the Medallion Lakehouse Architecture on Azure Data Lake Storage (ADLS Gen2).
- Designed dimensional data models using Star Schema to enable efficient downstream reporting and business analytics.
- Implemented Slowly Changing Dimension (SCD Type 2) logic for customer and product dimensions to preserve historical data integrity for sales and inventory reporting.
- Developed automated data quality validation frameworks including schema validation, row-count reconciliation, null checks, duplicate detection, and schema drift monitoring to improve data reliability.

PROJECTS

Swiggy AI Data Engineering Pipeline

[GitHub Repo](#)

Python, SQL, AWS S3, Snowflake, dbt, Apache Airflow, Docker, Ollama, RAG, Streamlit

- Built an end-to-end data engineering pipeline for a Swiggy-inspired food-delivery dataset, ingesting data from Amazon S3 to Snowflake and developing analytics-ready datasets using ELT practices.
- Implemented Medallion Architecture with dbt, developing modular staging models, Star Schema fact/dimension tables, data cleaning, standardization, and incremental transformations.
- Automated data workflows using Apache Airflow and Docker, orchestrating ingestion, transformation, and downstream analytics processing through repeatable pipelines.
- Applied SQL and Python-based data transformations to prepare reliable datasets for delivery trends, growth metrics, cancellations, operational performance, and customer analytics.
- Integrated RAG and Text-to-SQL with Ollama to enable natural-language querying of curated warehouse data and generate insights from operational metrics and customer reviews.
- Developed Streamlit and Looker Studio analytics dashboards for delivery trends, growth metrics, cancellations, delivery performance, and customer sentiment.

EDUCATION

Bachelor of Technology (B.Tech) in Computer Science

Graduated: 2023

Pranveer Singh Institute of Technology

Kanpur, Uttar Pradesh

- CGPA: 7.9 / 10

CERTIFICATIONS

- AWS Certified Solutions Architect – Associate**